

MR-API: A Comprehensive API Framework for Heterogeneous Multi-core Systems using Map Reduce Programming Model

Prashanth Thinakaran[†], Srinivas Avireddy[†], Prasanna Ranganathan[†], Sumalatha Ramachandran[‡]

[†]Undergraduate Students, Department of Information Technology, Anna University, Chennai-44

[‡]Associate Professor, Department of Information Technology, Anna University, Chennai-44

Motivation for an effective API

- ◆ Existing cloud systems map general purpose computations to CPU
- ◆ Difficulty in Automation of mapping computation to heterogeneous processing unit
- ◆ Abstraction in mapping instructions to heterogeneous unit
- ◆ Tapping GPU resources in cloud to improve the performance of general purpose applications
- ◆ Evolving effective heuristics for instruction delegation across heterogeneous cores
- ◆ Performance optimization of DAG's

MR-API Features

- ◆ Comprehensive Application Programming Interface for Heterogeneous Multi-core Multi threaded Systems
- ◆ Runtime aggregation of application specific instructions
- ◆ Enhances parallelism through processor level parallelism
- ◆ Mapping computations at runtime for heterogeneous processing units
- ◆ Provides scalability, fault tolerance and load balancing
- ◆ Performance optimization of DAG's using operation coalescing and removal of redundant temporary arrays
- ◆ Programming model of MR-API is library based improving the level of Abstraction
- ◆ Performance improvement by **50-90** percent in contrast to the serial CPU only architectures

Perks of MR-API

- ◆ Task based segregation Viz. Heuristics to handle flat and block data to achieve better granularity
- ◆ MR-API decreases the code complexity by hiding low level GPGPU constructs
- ◆ Application speedup increases at nearly exponential rate for N-body simulation which has $O(n^2)$ complexity

To integrate MR-API with Hadoop, the developer just needs to add two parameters to Hadoop's configuration file:

- ☞ `Djava.library.path= location of aparapi native library`
- ☞ `mapred.map.tasks=1`

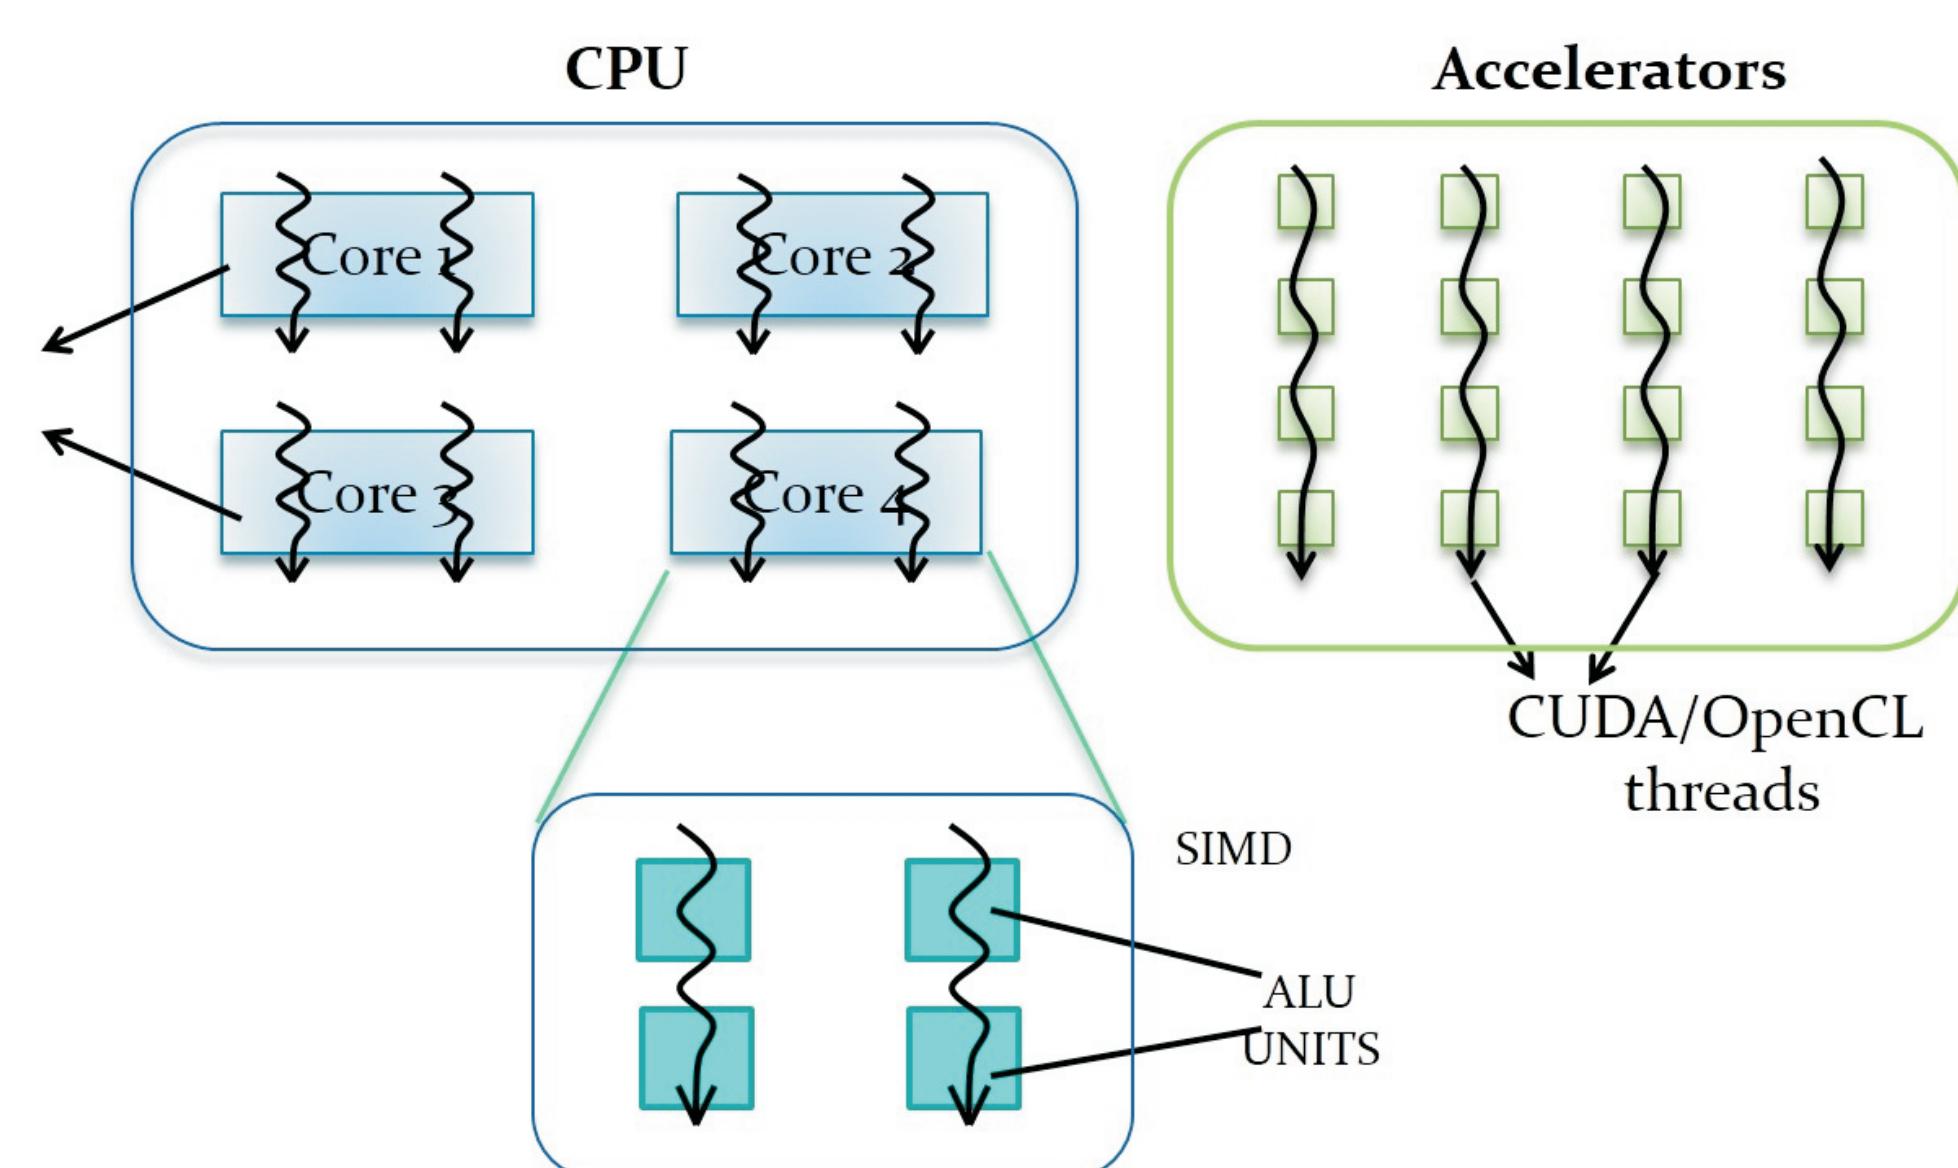


Fig 1: Typical Heterogeneous architecture comprising of CPU and GPUs

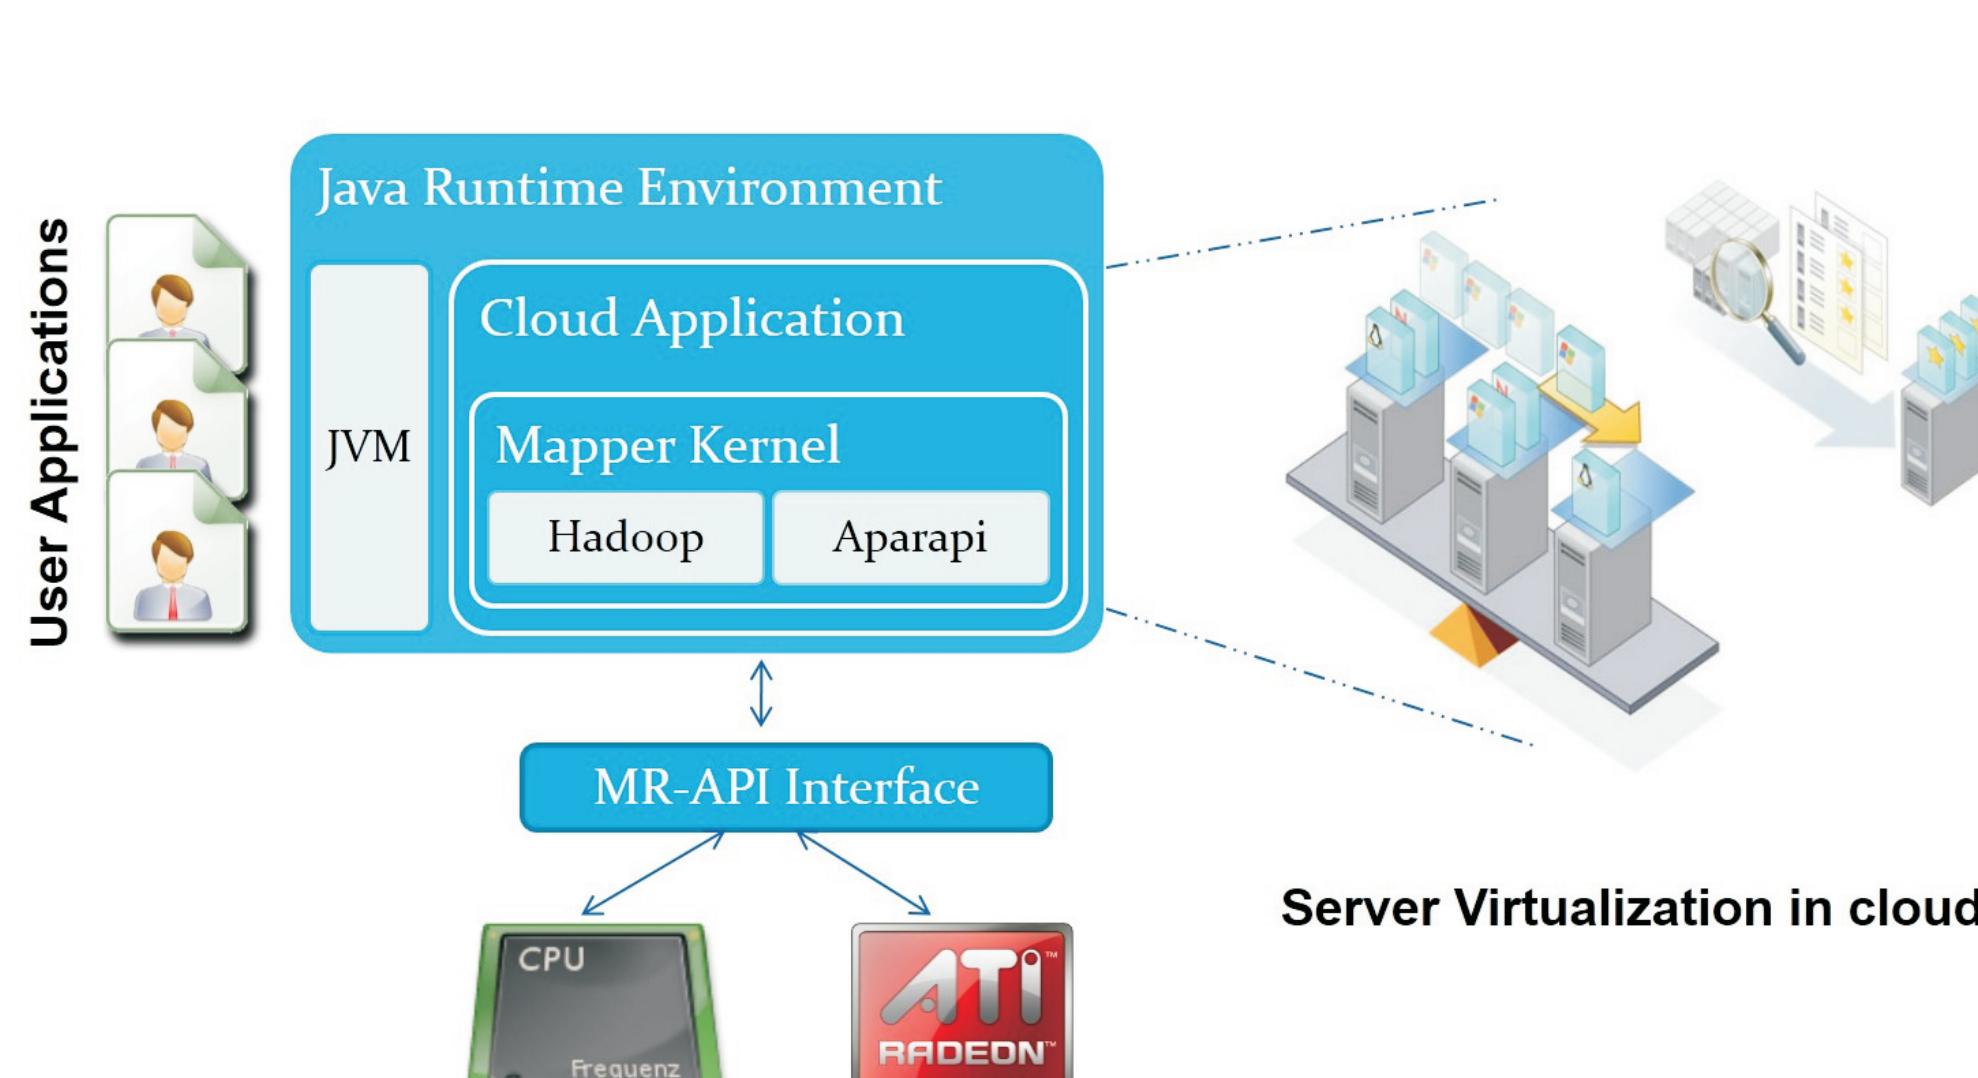


Fig 2: MR-API Software Architecture Framework

```
public class SimpleMapper<K1, V1, K2, V2> implements MapRunnable<K1, V1, K2, V2> {
    public void map(LongWritable offset, LongWritable size, OutputCollector<BooleanWritable, BooleanWritable> out, Reporter reporter) throws IOException {
        long numInside = 0L;
        long numOutside = 0L;

        final Point[] points = new Point[size.get()];
        for(int i=0;i<size.get();i++) {
            points[i] = new Point(Math.random(), Math.random());
            final float x = points[i].x * 0.5f;
            final float y = points[i].y * 0.5f;
            if(x+y > 0.25f)
                numOutside++;
            else
                numInside++;
        }

        out.collect(new BooleanWritable(true), new LongWritable(numInside));
        out.collect(new BooleanWritable(false), new LongWritable(numOutside));
    }

    public void run(RecordReader<K1, V1> input, OutputCollector<K2, V2> output, Reporter reporter) throws IOException {
        //allocate key&value instances
        //that are reused for all entries
        K1 key = input.createKey();
        V1 value = input.createValue();

        while(input.next(key, value)) {
            //map pair to output
            map((LongWritable)key, (LongWritable)value, (OutputCollector<BooleanWritable, LongWritable>)output, reporter);
        }
    }
}
```

Fig 4: Original MAP-REDUCE Program

```
1 public class GPUMapper extends
2 //the first two type are inputkey/valuepair
3 //the last two type are outputkey/valuepair
4 MapperKernel2FloatToBoolean<LongWritable, LongWritable,
5 BooleanWritable, LongWritable>{
6
7 @Override
8 public List<FloatTuple2> preprocess(
9 RecordReader<LongWritable, LongWritable> input,
10 Reporter reporter) throws IOException{
11
12 LongWritable key = input.createKey();
13 LongWritable value = input.createValue();
14 ArrayList<FloatTuple2> allGpuIn =
15 new ArrayList<FloatTuple2>();
16
17 while(input.next(key,value)){
18 for(int i=0;i<value.get();key.get();i++)
19 i++)
20 allGpuIn.add(new FloatTuple2(
21 (float) Math.random(),
22 (float) Math.random()));
23 }
24 return allGpuIn;
25 }
26
27
28 public boolean gpu(float x, float y){
29 return x + y > 0.25f;
30 }
31
32
33 @Override
34 public void postprocess(List<Boolean> gpuOut,
35 OutputCollector<BooleanWritable, LongWritable>
36 output)
37 throws IOException{
38 int numInside = 0;
39 int numOutside = 0;
40 for(boolean x:gpuOut)
41 if(x)
42 numInside++;
43 else
44 numOutside++;
45 output.collect(new BooleanWritable(true),
46 new LongWritable(numInside));
47 output.collect(new BooleanWritable(false),
48 new LongWritable(numOutside));
49 }
50 }
```

Fig 5: Modified MAP-REDUCE Program with MR-API Mapper Modules

Raw Instructions

DAG Builder

Categorization of GP instructions and accelerator specific instructions

Runtime MAP Reduce By spawning threads for CPU and Accelerators

Load balancing done by Scheduler

DAG Optimizer

Code Generation

Fig 3: Compilation sequence in MR-API based heterogeneous architectures

Sample Results

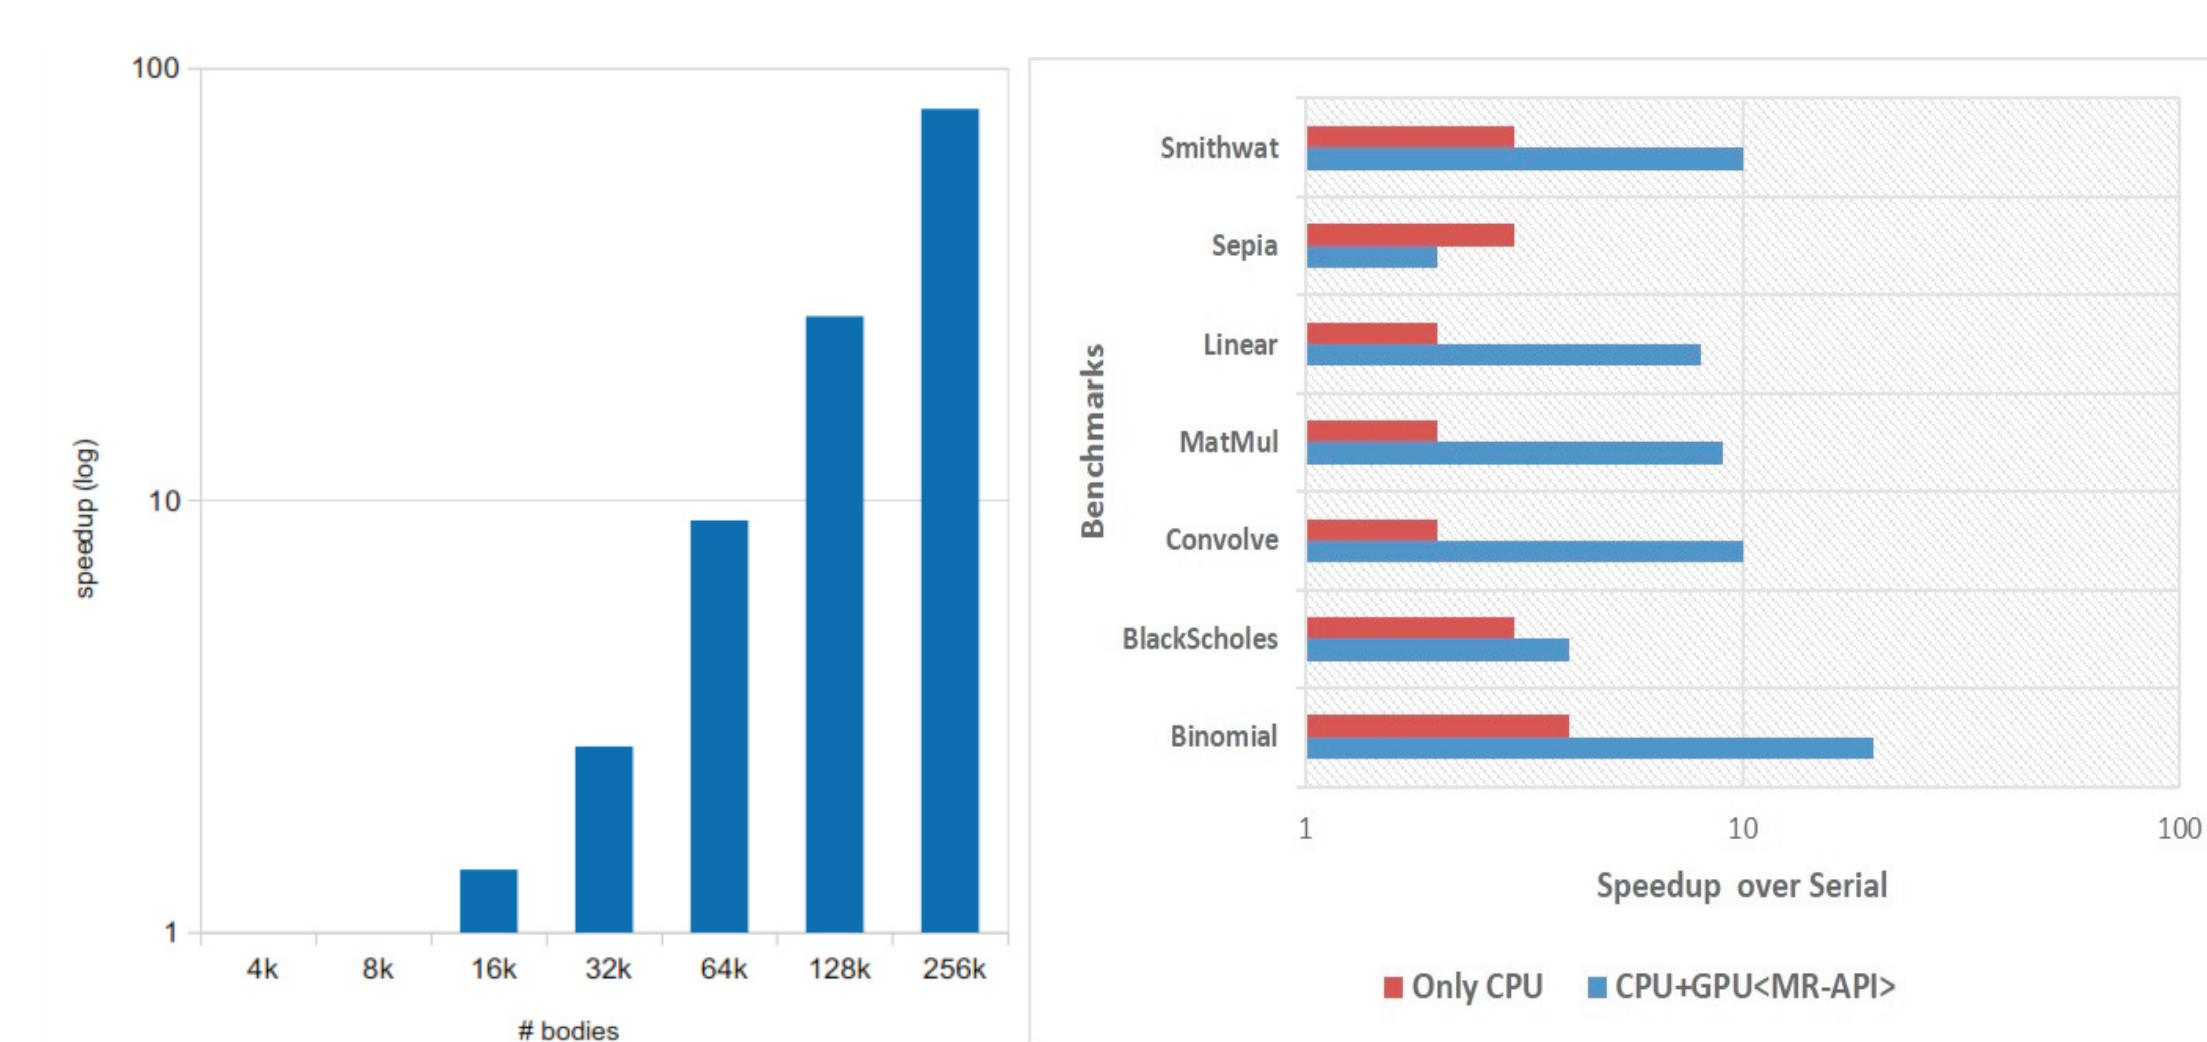


Fig 6: Speed up of Nbody simulation MR-API over CPU

Fig 7: Speed up of other benchmarks MR-API over CPU

Future work

- ◆ Evolving MR-API framework to support multiple languages apart from Java
- ◆ Developing libraries to support varied set of accelerators including **CUDA** powered Nvidia GPUs
- ◆ Complete automation of mapping framework to support hybrid execution without programmer intervention