

MR-API: A Comprehensive API Framework for Heterogeneous Multi-core Systems using Map Reduce Programming Model

Prashanth Thinakaran[†], Srinivas Avireddy[†], Prasanna Ranganathan[†], Sumalatha Ramachandran[†]

[†]Undergraduate Students, Department of Information Technology, Anna University, Chennai-44

[†]Associate Professor, Department of Information Technology, Anna University, Chennai-44

Motivation for an effective API

- ◆ Existing cloud systems map general purpose computations to CPU
- ◆ Difficulty in Automation of mapping computation to heterogeneous processing unit
- ◆ Abstraction in mapping instructions to heterogeneous unit
- ◆ Tapping GPU resources in cloud to improve the performance of general purpose applications
- ◆ Evolving effective heuristics for instruction delegation across heterogeneous cores
- ◆ Performance optimization of DAG's

MR-API Features

- ◆ Comprehensive Application Programming Interface for Heterogeneous Multi-core Multi-threaded Systems
- ◆ Runtime aggregation of application specific instructions
- ◆ Enhances parallelism through processor level parallelism
- ◆ Mapping computations at runtime for heterogeneous processing units
- ◆ Provides scalability, fault tolerance and load balancing
- ◆ Performance optimization of DAG's using operation coalescing and removal of redundant temporary arrays
- ◆ Programming model of **MR-API** is library based improving the level of Abstraction
- ◆ Performance improvement by **50-90** percent in contrast to the serial CPU only architectures

Perks of MR-API

- ◆ Task based segregation Viz. Heuristics to handle flat and block data to achieve better granularity
- ◆ MR-API decreases the code complexity by hiding low level GPGPU constructs
- ◆ Application speedup increases at nearly exponential rate for N-body simulation which has $O(n^2)$ complexity

To integrate MR-API with Hadoop, the developer just needs to add two parameters to Hadoop's configuration file:

- ☞ `Djava.library.path= location of aparapi native library`
- ☞ `mapred.map.tasks=1`

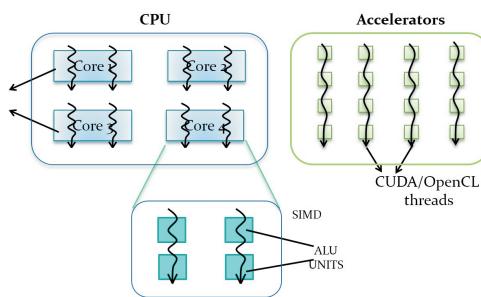


Fig 1: Typical Heterogeneous architecture comprising of CPU and GPUs

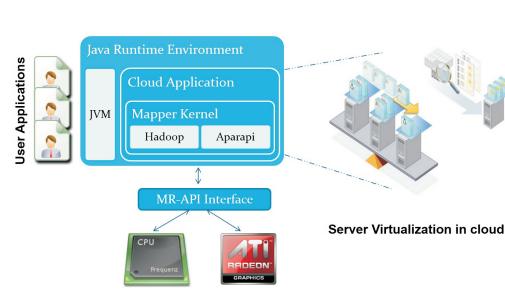


Fig 2: MR-API Software Architecture Framework

```
public class Simple Mapper<K1 , V1 , K2 , V2>
1 implements MapWritable<K1 , V1 , K2 , V2> f
3
4 public void map ( LongWritable offset ,
5   LongWritable size ,
6   OutputCollector< Boolean Writable ,
7   LongWritable > out ,
8   Reporter reporter ) throws IOException
9   long numInside = 0L ;
10  long numOutside = 0L ;
11
12 final Point[] points =
13  new Point [ size . get () ] ;
14
15 for ( int i = 0 ; i < size . get () ; i ++ ) f
16  points [ i ] = new Point ( Math . random () ,
17  Math . random () ) ;
18  final float x = points [ i ] . x * 0.5f ;
19  final float y = points [ i ] . y * 0.5f ;
20  if ( x + y > 0.25f )
21    numInside ++ ;
22  else {
23    numOutside ++ ;
24  }
25}
26
27 out . collect ( new Boolean Writable ( true ) ,
28  new LongWritable ( numInside ) ) ;
29 out . collect ( new Boolean Writable ( false ) ,
30  new LongWritable ( numOutside ) ) ;
31}
32
33 public void run ( RecordReader< K1 , V1 > input ,
34  OutputCollector< K2 , V2 > output ,
35  Reporter reporter ) throws IOException
36
37 // allocate key & value instances
38 // that are reused for all entries
39 K1 key = input . createKey () ;
40 V1 value = input . createValue () ;
41
42 while ( input . next ( key , value ) ) f
43  // map pair to output
44  map ( ( LongWritable ) key ,
45  ( LongWritable ) value ,
46  ( OutputCollector< Boolean Writable ,
47  LongWritable > ) output , reporter ) ;
48 }
49 }
50}
```

Fig 4: Original MAP-REDUCE Program

```
1 public class GPUMapper extends
2 // the first two type are inputkey/valuepair
3 // the last two type are outputkey/valuepair
4 Mapper Kernel< FloatToBoolean , LongWritable , LongWritable ,
5 Boolean Writable , LongWritable > i
6
7 @Override
8 public List< Float Tuple2 > preprocess (
9  RecordReader< LongWritable , LongWritable > input ,
10 Reporter reporter ) throws IOException {
11
12  LongWritable key = input . createKey () ;
13  LongWritable value = input . createValue () ;
14  ArrayList< Float Tuple2 > allGpuIn
15  = new ArrayList< Float Tuple2 > ( ) ;
16
17 while ( input . next ( key , value ) ) {
18 for ( int i = 0 ; i < value . get () . key . get () ;
19  i ++ )
20  allGpuIn . add ( new Float Tuple2 (
21  ( float ) Math . random () ,
22  ( float ) Math . random () ) ) ;
23  }
24  return allGpuIn ;
25 }
26
27
28 public boolean gpu ( float x , float y ) {
29 return x + y > 0.25f ;
30 }
31
32
33 @Override
34 public void postprocess ( List< Boolean > gpuOut ,
35  OutputCollector< Boolean Writable , LongWritable >
36  output )
37 throws IOException {
38  int numInside = 0 ;
39  int numOutside = 0 ;
40  for ( boolean x : gpuOut ) {
41    if ( x )
42      numInside ++ ;
43    else
44      numOutside ++ ;
45  output . collect ( new Boolean Writable ( true ) ,
46  new LongWritable ( numInside ) ) ;
47  output . collect ( new Boolean Writable ( false ) ,
48  new LongWritable ( numOutside ) ) ;
49  }
50 }
```

Fig 5: Modified MAP-REDUCE Program with MR-API Mapper Modules

Raw Instructions

↓

DAG Builder

↓

Categorization of GP instructions and accelerator specific instructions

↓

Runtime MAP Reduce By spawning threads for CPU and Accelerators

↓

Load balancing done by Scheduler

↓

DAG Optimizer

↓

Code Generation

Fig 3: Compilation sequence in MR-API based heterogeneous architectures

Sample Results

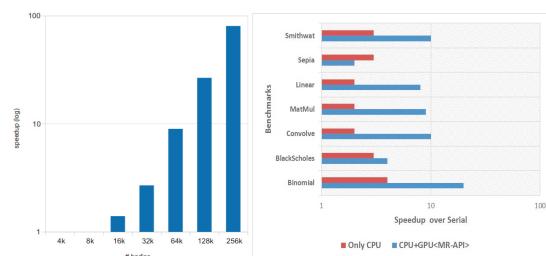


Fig 6: Speed up of Nbody simulation MR-API over CPU

Fig 7: Speed up of other benchmarks MR-API over CPU

Future work

- ◆ Evolving MR-API framework to support multiple languages apart from Java
- ◆ Developing libraries to support varied set of accelerators including **CUDA** powered Nvidia GPUs
- ◆ Complete automation of mapping framework to support hybrid execution without programmer intervention